# Package 'kamila'

**Type** Package

**Version** 0.1.1.1

**Date** 2016-08-18

**Author** Alexander Foss [aut, cre],
  Marianthi Markatou [aut]

**Maintainer** Alexander Foss <alexanderhfoss@gmail.com>

**Title** Methods for Clustering Mixed-Type Data

**Description** Implements methods for clustering mixed-type data,
  specifically combinations of continuous and nominal data. Special attention
  is paid to the often-overlooked problem of equitably balancing the
  contribution of the continuous and categorical variables. This package
  implements KAMILA clustering, a novel method for clustering
  mixed-type data in the spirit of k-means clustering. It does not require
  dummy coding of variables, and is efficient enough to scale to rather large
  data sets. Also implemented is Modha-Spangler clustering, which uses a
  brute-force strategy to maximize the cluster separation simultaneously in the
  continuous and categorical variables.

**Depends** R (>= 3.0.0)

**License** GPL-3

**URL** https://github.com/ahfoss/kamila

**BugReports** https://github.com/ahfoss/kamila/issues

**Imports** stats, abind, KernSmooth, gtools, Rcpp, mclust, plyr

**LinkingTo** Rcpp

**Suggests** testthat, clustMD, ggplot2, Hmisc

**RoxygenNote** 5.0.1

**NeedsCompilation** yes

**Repository** CRAN

**Date/Publication** 2016-08-19 00:46:49

# R topics documented:

---

kamila-package | *Clustering for mixed continuous and categorical data sets*

---

#### Description

A collection of methods for clustering mixed type data, including KAMILA (KAy-means for MIxed LArge data) and a flexible implementation of Modha-Spangler clustering

#### Details

| | |
|---|---|
| Package: | kamila |
| Type: | Package |
| Version: | 0.1.0 |
| Date: | 2015-10-06 |
| License: | GPL-3 |

#### Author(s)

Alex Foss and Marianthi Markatou

Maintainer: Alex Foss <alexanderhfoss@gmail.com>

#### References

AH Foss, M Markatou, B Ray, and A Heching (in press). A semiparametric method for clustering mixed data. Machine Learning, DOI: 10.1007/s10994-016-5575-7.

DS Modha and S Spangler (2003). Feature weighting in k-means clustering. Machine Learning 52(3), 217-237.

#### Examples

```
## Not run:
# import and format a mixed-type data set
```

```
data(Byar, package='clustMD')
Byar$logSpap <- log(Byar$Serum.prostatic.acid.phosphatase)
conInd <- c(5,6,8:10,16)
conVars <- Byar[,conInd]
conVars <- data.frame(scale(conVars))

catVarsFac <- Byar[,-c(1:2,conInd,11,14,15)]
catVarsFac[] <- lapply(catVarsFac, factor)
catVarsDum <- dummyCodeFactorDf(catVarsFac)

# Modha-Spangler clustering with kmeans default Hartigan-Wong algorithm
gmsResHw <- gmsClust(conVars, catVarsDum, nclust = 3)

# Modha-Spangler clustering with kmeans Forgy-Lloyd algorithm
# NOTE searchDensity should be >= 10 for optimal performance:
# this is just a syntax demo
gmsResLloyd <- gmsClust(conVars, catVarsDum, nclust = 3,
  algorithm = "Lloyd", searchDensity = 3)

# KAMILA clustering
kamRes <- kamila(conVars, catVarsFac, numClust=3, numInit=10)

# Plot results
ternarySurvival <- factor(Byar$SurvStat)
levels(ternarySurvival) <- c('Alive','DeadProst','DeadOther')[c(1,2,rep(3,8))]
plottingData <- cbind(
  conVars,
  catVarsFac,
  KamilaCluster = factor(kamRes$finalMemb),
  MSCluster = factor(gmsResHw$results$cluster))
plottingData$Bone.metastases <- ifelse(
  plottingData$Bone.metastases == '1', yes='Yes',no='No')

# Plot Modha-Spangler/Hartigan-Wong results
msPlot <- ggplot(
  plottingData,
  aes(
    x=logSpap,
    y=Index.of.tumour.stage.and.histolic.grade,
    color=ternarySurvival,
    shape=MSCluster))
plotOpts <- function(pl) (pl + geom_point() +
  scale_shape_manual(values=c(2,3,7)) + geom_jitter())
plotOpts(msPlot)

# Plot KAMILA results
kamPlot <- ggplot(
  plottingData,
  aes(
    x=logSpap,
    y=Index.of.tumour.stage.and.histolic.grade,
    color=ternarySurvival,
    shape=KamilaCluster))
```

```
plotOpts(kamPlot)

## End(Not run)
```

---

classifyKamila          *Classify new data into existing KAMILA clusters*

---

### Description

A function that classifies a new data set into existing KAMILA clusters using the output object from the kamila function.

### Usage

```
classifyKamila(obj, newData)
```

### Arguments

| | |
|---|---|
| obj | An output object from the kamila function. |
| newData | A list of length 2, with first element a data frame of continuous variables, and second element a data frame of categorical factors. |

### Details

A function that takes obj, the output from the kamila function, and newData, a list of length 2, where the first element is a data frame of continuous variables, and the second element is a data frame of categorical factors. Both data frames must have the same format as the original data used to construct the kamila clustering.

### Value

An integer vector denoting cluster assignments of the new data points.

---

dummyCodeFactorDf          *Dummy coding of a data frame of factor variables*

---

### Description

Given a data frame of factor variables, this function returns a numeric matrix of 0–1 dummy-coded variables.

### Usage

```
dummyCodeFactorDf(dat)
```

## Arguments

| | |
|---|---|
| dat | A data frame of factor variables |

## Value

A numeric matrix of 0–1 dummy coded variables

## Examples

```
dd <- data.frame(a=factor(1:8), b=factor(letters[1:8]))
dummyCodeFactorDf(dd)
```

---

| genMixedData | *Generate simulated mixed-type data with cluster structure.* |
|---|---|

---

## Description

This function simulates mixed-type data sets with a latent cluster structure, with continuous and nominal variables.

## Usage

```
genMixedData(sampSize, nConVar, nCatVar, nCatLevels, nConWithErr, nCatWithErr,
  popProportions, conErrLev, catErrLev)
```

## Arguments

| | |
|---|---|
| sampSize | Integer: Size of the simulated data set. |
| nConVar | The number of continuous variables. |
| nCatVar | The number of categorical variables. |
| nCatLevels | Integer: The number of categories per categorical variables. Currently must be a multiple of the number of populations specified in popProportions. |
| nConWithErr | Integer: The number of continuous variables with error. |
| nCatWithErr | Integer: The number of categorical variables with error. |
| popProportions | A vector of scalars that sums to one. The length gives the number of populations (clusters), with values denoting the prior probability of observing a member of the corresponding population. NOTE: currently only two populations are supported. |
| conErrLev | A scalar between 0.01 and 1 denoting the univariate overlap between clusters on the continuous variables specified to have error. |
| catErrLev | Univariate overlap level for the categorical variables with error. |

**Details**

This function simulates mixed-type data sets with a latent cluster structure. Continuous variables follow a normal mixture model, and categorical variables follow a multinomial mixture model. Overlap of the continuous and categorical variables (i.e. how clear the cluster structure is) can be manipulated by the user. The default overlap level is 1 percent (i.e. almost perfect separation), and a user-specified number of continuous and categorical variables can be specified to be measured with error, in which case the overlap can be selectively set to be anywhere within 1 and 100 percent (100 percent corresponds to complete overlap).

NOTE: Currently, only two populations (clusters) are supported.

**Value**

A list with the following elements:

| | |
|---|---|
| trueID | Integer vector giving population (cluster) membership of each observation |
| trueMus | Mean parameters used for population (cluster) centers in the continuous variables |
| conVars | The continuous variables |
| errVariance | Variance parameter used for continuous error distribution |
| popProbsNoErr | Multinomial probability vectors for categorical variables without measurement error |
| popProbsWithErr | |
| | Multinomial probability vectors for categorical variables with measurement error |
| catVars | The categorical variables |

**Examples**

```
dat <- genMixedData(100, 2, 2, nCatLevels=4, nConWithErr=1, nCatWithErr=1,
  popProportions=c(0.3,0.7), conErrLev=0.3, catErrLev=0.2)
with(dat,plot(conVars,col=trueID))
with(dat,table(data.frame(catVars[,1:2],trueID)))
```

---

| gmsClust | *A general implementation of Modha-Spangler clustering for mixed-type data.* |
|---|---|

---

**Description**

Modha-Spangler clustering estimates the optimal weighting for continuous vs categorical variables using a brute-force search strategy.

**Usage**

```
gmsClust(conData, catData, nclust, searchDensity = 10, clustFun = wkmeans,
  conDist = squaredEuc, catDist = squaredEuc, ...)
```

## Arguments

| | |
|---|---|
| conData | A data frame of continuous variables. |
| catData | A data frame of categorical variables; the allowable variable types depend on the specific clustering function used. |
| nclust | An integer specifying the number of clusters. |
| searchDensity | An integer determining the number of distinct cluster weightings evaluated in the brute-force search. |
| clustFun | The clustering function to be applied. |
| conDist | The continuous distance function used to construct the objective function. |
| catDist | The categorical distance function used to construct the objective function. |
| ... | Arguments to be passed to the clustFun. |

## Details

Modha-Spangler clustering uses a brute-force search strategy to estimate the optimal weighting for continuous vs categorical variables. This implementation admits an arbitrary clustering function and arbitrary objective functions for continuous and categorical variables.

The input parameter clustFun must be a function accepting inputs (conData, catData, conWeight, nclust, ...) and returning a list containing (at least) the elements cluster, conCenters, and catCenters. The list element "cluster" contains cluster memberships denoted by the integers 1:nclust. The list elements "conCenters" and "catCenters" must be data frames whose rows denote cluster centroids. The function clustFun must allow nclust = 1, in which case $centers returns a data frame with a single row. Input parameters conDist and catDist are functions that must each take two data frame rows as input and return a scalar distance measure.

## Value

A list containing the following results objects:

| | |
|---|---|
| results | A results object corresponding to the base clustering algorithm |
| objFun | A numeric vector of length searchDensity containing the values of the objective function for each weight used |
| Qcon | A numeric vector of length searchDensity containing the values of the continuous component of the objective function |
| Qcon | A numeric vector of length searchDensity containing the values of the categorical component of the objective function |
| bestInd | The index of the most successful run |
| weights | A numeric vector of length searchDensity containing the continuous weights used |

## Examples

```
## Not run:
# Generate toy data set with poor quality categorical variables and good
# quality continuous variables.
```

```
set.seed(1)
dat <- genMixedData(200, nConVar=2, nCatVar=2, nCatLevels=4, nConWithErr=2,
  nCatWithErr=2, popProportions=c(.5,.5), conErrLev=0.3, catErrLev=0.8)
catDf <- dummyCodeFactorDf(data.frame(apply(dat$catVars, 2, factor)))
conDf <- data.frame(scale(dat$conVars))

msRes <- gmsClust(conDf, catDf, nclust=2)

table(msRes$results$cluster, dat$trueID)

## End(Not run)
```

---

kamila                          *KAMILA clustering of mixed-type data.*

---

## Description

KAMILA is an iterative clustering method that equitably balances the contribution of continuous and categorical variables.

## Usage

```
kamila(conVar, catFactor, numClust, numInit, conWeights = rep(1,
  ncol(conVar)), catWeights = rep(1, ncol(catFactor)), maxIter = 25,
  conInitMethod = "runif", catBw = 0.025, verbose = FALSE,
  calcNumClust = "none", numPredStrCvRun = 10, predStrThresh = 0.8)
```

## Arguments

| | |
|---|---|
| conVar | A data frame of continuous variables. |
| catFactor | A data frame of factors. |
| numClust | The number of clusters returned by the algorithm. |
| numInit | The number of initializations used. |
| conWeights | A vector of continuous weights for the continuous variables. |
| catWeights | A vector of continuous weights for the categorical variables. |
| maxIter | The maximum number of iterations in each run. |
| conInitMethod | Character: The method used to initialize each run. |
| catBw | The bandwidth used for the categorical kernel. |
| verbose | Logical: Whether detailed results should be printed and returned. |
| calcNumClust | Character: Method for selecting the number of clusters. |
| numPredStrCvRun | |
| | Numeric: Number of CV runs for prediction strength method. Ignored unless calcNumClust == 'ps' |
| predStrThresh | Numeric: Threshold for prediction strength method. Ignored unless calcNumClust == 'ps' |

## Details

KAMILA (KAy-means for MIxed LArge data sets) is an iterative clustering method that equitably balances the contribution of the continuous and categorical variables. It uses a kernel density estimation technique to flexibly model spherical clusters in the continuous domain, and uses a multinomial model in the categorical domain.

Weighting scheme: If no weights are desired, set all weights to 1 (the default setting). Let a_1, ..., a_p denote the weights for p continuous variables. Let b_1, ..., b_q denote the weights for q categorical variables. Currently, continuous weights are applied during the calculation of Euclidean distance, as: Categorical weights are applied to the log-likelihoods obtained by the level probabilities given cluster membership as: Total log likelihood for the kth cluster is obtained by weighting the single continuous log-likelihood by the mean of all continuous weights plus logLikCat_k: Note that weights between 0 and 1 are admissible; weights equal to zero completely remove a variable's influence on the clustering; weights equal to 1 leave a variable's contribution unchanged. Weights between 0 and 1 may not be comparable across continuous and categorical variables. Estimating the number of clusters: Default is no estimation method. Setting calcNumClust to 'ps' uses the prediction strength method of Tibshirani & Walther (J. of Comp. and Graphical Stats. 14(3), 2005). There is no perfect method for estimating the number of clusters; PS tends to give a smaller number than, say, BIC based methods for large sample sizes. The user must specify the number of cross-validation runs and the threshold for determining the number of clusters. The smaller the threshold, the larger the number of clusters selected.

## Value

A list with the following results objects:

| | |
|---|---|
| `finalMemb` | A numeric vector with cluster assignment indicated by integer. |
| `numIter` | |
| `finalLogLik` | The pseudo log-likelihood of the returned clustering. |
| `finalObj` | |
| `finalCenters` | |
| `finalProbs` | |
| `input` | Object with the given input parameter values. |
| `nClust` | An object describing the results of selecting the number of clusters, empty if calcNumClust == 'none'. |
| `verbose` | An optionally returned object with more detailed information. |

## Examples

```
# Generate toy data set with poor quality categorical variables and good
# quality continuous variables.
set.seed(1)
dat <- genMixedData(200, nConVar=2, nCatVar=2, nCatLevels=4, nConWithErr=2,
  nCatWithErr=2, popProportions=c(.5,.5), conErrLev=0.3, catErrLev=0.8)
catDf <- data.frame(apply(dat$catVars, 2, factor))
conDf <- data.frame(scale(dat$conVars))

kamRes <- kamila(conDf, catDf, numClust=2, numInit=10)
```

```
table(kamRes$finalMemb, dat$trueID)
```

---

wkmeans                              *Weighted k-means for mixed-type data*

---

### Description

Weighted k-means for mixed continuous and categorical variables. A user-specified weight `conWeight` controls the relative contribution of the variable types to the cluster solution.

### Usage

```
wkmeans(conData, catData, conWeight, nclust, ...)
```

### Arguments

| | |
|---|---|
| conData | The continuous variables. Must be coercible to a data frame. |
| catData | The categorical variables, either as factors or dummy-coded variables. Must be coercible to a data frame. |
| conWeight | The continuous weight; must be between 0 and 1. The categorical weight is `1-conWeight`. |
| nclust | The number of clusters. |
| ... | Optional arguments passed to `kmeans`. |

### Details

A simple adaptation of `stats::kmeans` to mixed-type data. Continuous variables are multiplied by the input parameter `conWeight`, and categorical variables are multipled by `1-conWeight`. If factor variables are input to `catData`, they are transformed to 0-1 dummy coded variables with the function dummyCodeFactorDf.

### Value

A stats::kmeans results object, with additional slots `conCenters` and `catCenters` giving the actual centers adjusted for the weighting process.

### See Also

[dummyCodeFactorDf](#)

[kmeans](#)

## Examples

```
# Generate toy data set with poor quality categorical variables and good
# quality continuous variables.
set.seed(1)
dat <- genMixedData(200, nConVar=2, nCatVar=2, nCatLevels=4, nConWithErr=2,
  nCatWithErr=2, popProportions=c(.5,.5), conErrLev=0.3, catErrLev=0.8)
catDf <- data.frame(apply(dat$catVars, 2, factor))
conDf <- data.frame(scale(dat$conVars))

# A clustering that emphasizes the continuous variables
r1 <- with(dat,wkmeans(conDf, catDf, 0.9, 2))
table(r1$cluster, dat$trueID)

# A clustering that emphasizes the categorical variables; note argument
# passed to the underlying stats::kmeans function
r2 <- with(dat,wkmeans(conDf, catDf, 0.1, 2, nstart=4))
table(r2$cluster, dat$trueID)
```

# Index